



# Introduction to Pandas

A Library that is Used for Data  
Manipulation and Analysis Tool

Using Powerful Data Structures

# Pandas First Steps: **install** and **import**

- ▶ Pandas is an easy package to install. Open up your terminal program (shell or cmd) and install it using either of the following commands:
- ▶ For **jupyter notebook** users, you can run this cell:
  - ▶ The **!** at the beginning runs cells as if they were in a terminal.
- ▶ To import pandas we usually import it with a shorter name since it's used so much:


```
$ conda install pandas  
OR  
$ pip install pandas
```

```
!pip install pandas
```

```
import pandas as pd
```

# pandas: Data Table Representation

DataFrame



# Core components of pandas:

## Series & DataFrames

- ▶ The primary two components of pandas are the Series and DataFrame.
  - ▶ **Series** is essentially a **column**, and
  - ▶ **DataFrame** is a multi-dimensional table made up of a **collection of Series**.
- ▶ **DataFrames** and **Series** are quite similar in that many operations that you can do with one you can do with the other, such as filling in null values and calculating the mean.
  - ▶ A Data frame is a two-dimensional data structure, i.e., data is aligned in a tabular fashion in rows and columns.

- **Features of DataFrame**

- Potentially columns are of different types
- Size - Mutable
- Labeled axes (*rows* and *columns*)
- Can Perform Arithmetic operations on rows and columns

Series

	apples
0	3
1	2
2	0
3	1

+

Series

	oranges
0	0
1	3
2	7
3	2

=

DataFrame

	apples	oranges
0	3	0
1	2	3
2	0	7
3	1	2

rows

# Types of Data Structure in Pandas

Data Structure	Dimensions	Description
Series	1	1D labeled <u>homogeneous</u> array with immutable size
Data Frames	2	General 2D labeled, size mutable tabular structure with potentially <u>heterogeneously</u> typed columns.
Panel	3	General 3D labeled, size mutable array.

- **Series & DataFrame**

- **Series** is a one-dimensional array (1D Array) like structure with homogeneous data.
- **DataFrame** is a two-dimensional array (2D Array) with heterogeneous data.

- **Panel**

- Panel is a three-dimensional data structure (3D Array) with heterogeneous data.
- It is hard to represent the panel in graphical representation.
- But a panel can be illustrated as a container of DataFrame

# pandas.DataFrame

```
pandas.DataFrame(data, index, columns, dtype, copy)
```

- ▶ **data:** data takes various forms like *ndarray*, *series*, *map*, *lists*, *dict*, constants and also another *DataFrame*.
- ▶ **index:** For the row labels, that are to be used for the resulting frame, Optional, Default is *np.arange(n)* if no index is passed.
- ▶ **columns:** For column labels, the optional default syntax is - *np.arange(n)*. This is only true if no index is passed.
- ▶ **dtype:** Data type of each column.
- ▶ **copy:** This command (or whatever it is) is used for copying of data, if the default is False.
- **Create DataFrame**
  - A pandas DataFrame can be created using various inputs like –
    - Lists
    - dict
    - Series
    - Numpy ndarrays
    - Another DataFrame





# Creating a DataFrame from scratch

# Creating a DataFrame from scratch

- ▶ There are many ways to create a DataFrame from scratch, but a great option is to just use a simple dict. But first you must import pandas.

```
import pandas as pd
```

- ▶ Let's say we have a fruit stand that sells apples and oranges. We want to have a column for each fruit and a row for each customer purchase. To organize this as a dictionary for pandas we could do something like:

```
data = { 'apples': [3, 2, 0, 1] , 'oranges': [0, 3, 7, 2] }
```

- ▶ And then pass it to the pandas DataFrame constructor:

```
df = pd.DataFrame(data)
```



	apples	oranges
0	3	0
1	2	3
2	0	7
3	1	2



# How did that work?

- ▶ Each (key, value) item in data corresponds to a **column** in the resulting **DataFrame**.
- ▶ The **Index** of this **DataFrame** was given to us on creation as the numbers 0–3, but we could also create our own when we initialize the **DataFrame**.
- ▶ E.g. if you want to have customer names as the **index**:

```
df = pd.DataFrame(data, index=['Ahmad', 'Ali', 'Rashed', 'Hamza'])
```

	apples	oranges
Ahmad	3	0
Ali	2	3
Rashed	0	7
Hamza	1	2

- So now we could locate a customer's order by using their names:

```
df.loc['Ali']
```

```
apples    2  
oranges    3  
Name: Ali, dtype: int64
```

# pandas.DataFrame.from\_dict

```
pandas.DataFrame.from_dict(data, orient='columns', dtype=None, columns=None)
```

- **data** : dict
  - Of the form {**field**:array-like} or {**field**:dict}.
- **orient** : { 'columns' , 'index' }, default 'columns'
  - The “orientation” of the data.
  - If the keys of the passed dict should be the columns of the resulting DataFrame, pass 'columns' (default).
  - Otherwise if the keys should be rows, pass 'index'.
- **dtype** : dtype, default None
  - Data type to force, otherwise infer.
- **columns** : list, default None
  - Column labels to use when **orient**='index'. Raises a **ValueError** if used with **orient**='columns'.

# pandas' **orient** keyword

```
data = {'col_1':[3, 2, 1, 0], 'col_2':['a','b','c','d']}  
pd.DataFrame.from_dict(data)
```



	col_1	col_2
0	3	a
1	2	b
2	1	c
3	0	d

```
data = {'row_1':[3, 2, 1, 0], 'row_2':['a','b','c','d']}  
pd.DataFrame.from_dict(data,  
                        orient='index')
```




	0	1	2	3
row_1	3	2	1	0
row_2	a	b	c	d

```
data = {'row_1':[3, 2, 1, 0], 'row_2':['a','b','c','d']}  
pd.DataFrame.from_dict(data,  
                        orient = 'index',  
                        columns = ['A','B','C','D'])
```

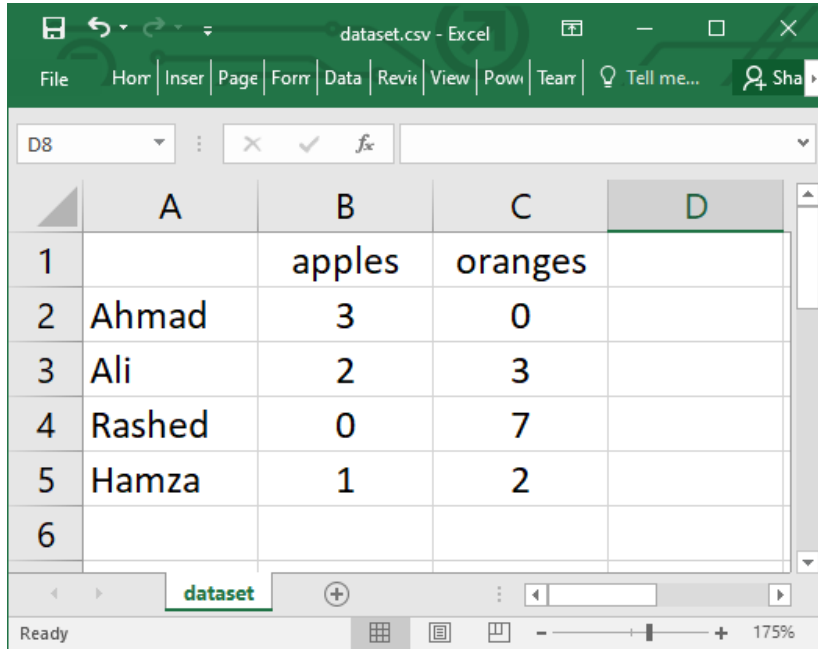


	A	B	C	D
row_1	3	2	1	0
row_2	a	b	c	d

# Loading a DataFrame from files



# Reading data from a CSV file



	A	B	C	D
1		apples	oranges	
2	Ahmad	3	0	
3	Ali	2	3	
4	Rashed	0	7	
5	Hamza	1	2	
6				



```
File Edit Format Run Options Window Help
1 import pandas as pd
2
3 df = pd.read_csv('dataset.csv')
4 print(df)
5
6 # OR
7
8 df = pd.read_csv('dataset.csv', index_col=0)
9 print(df)
10
```

Ln: 6 Col: 0

# Reading data from CSVs

- ▶ With CSV files, all you need is a single line to load in the data:

```
df = pd.read_csv('dataset.csv')
```

	Unnamed: 0	apples	oranges
0	Ahmad	3	0
1	Ali	2	3
2	Rashed	0	7
3	Hamza	1	2

- CSVs don't have indexes like our DataFrames, so all we need to do is just designate the `index_col` when reading:

```
df = pd.read_csv('dataset.csv', index_col=0)
```

- *Note: here we're setting the index to be column zero.*

	apples	oranges
Ahmad	3	0
Ali	2	3
Rashed	0	7
Hamza	1	2



# Reading data from JSON

- ▶ If you have a JSON file — which is essentially a stored Python dict — pandas can read this just as easily:

```
df = pd.read_json('dataset.json')
```

- Notice this time our index came with us correctly since using JSON allowed indexes to work through nesting.
- Pandas will try to figure out how to create a DataFrame by analyzing structure of your JSON, and sometimes it doesn't get it right.
- Often you'll need to set the `orient` keyword argument depending on the structure

# Example #1: Reading data from JSON

```
{  
  "apples" : {"Ahmad":3, "Ali":2, "Rashed":0, "Hamza":1},  
  "oranges": {"Ahmad":0, "Ali":3, "Rashed":7, "Hamza":2}  
}
```



```
File Edit Format Run Options Window Help  
1 import pandas as pd  
2  
3 df = pd.read_json('dataset.json')  
4 print(df)  
Ln: 1 Col: 0
```



	apples	oranges
Ahmad	3	0
Ali	2	3
Rashed	0	7
Hamza	1	2

## Example #2: Reading data from JSON

```
{
  "Ahmad" : {"apples":3, "oranges":0},
  "Ali" : {"apples":2, "oranges":3},
  "Rashed" : {"apples":0, "oranges":7},
  "Hamza" : {"apples":1, "oranges":2}
}
```



```
File Edit Format Run Options Window Help
1 import pandas as pd
2
3 df = pd.read_json('dataset.json')
4 print(df)
Ln: 1 Col: 0
```



	Ahmad	Ali	Rashed	Hamza
apples	3	2	0	1
oranges	0	3	7	2

# Example #3: Reading data from JSON

```
{
  "Ahmad" : {"apples":3, "oranges":0},
  "Ali" : {"apples":2, "oranges":3},
  "Rashed" : {"apples":0, "oranges":7},
  "Hamza" : {"apples":1, "oranges":2}
}
```

```
File Edit Format Run Options Window Help
1 import pandas as pd
2
3 df = pd.read_json('dataset.json',
4                   orient='column')
5 print(df)
```

Ln: 6 Col: 0

	Ahmad	Ali	Rashed	Hamza
apples	3	2	0	1
oranges	0	3	7	2

```
File Edit Format Run Options Window Help
1 import pandas as pd
2
3 df = pd.read_json('dataset.json',
4                   orient='index')
5 print(df)
```

Ln: 6 Col: 0

	apples	oranges
Ahmad	3	0
Ali	2	3
Rashed	0	7
Hamza	1	2

# Converting back to a CSV or JSON

- So after extensive work on cleaning your data, you're now ready to save it as a file of your choice. Similar to the ways we read in data, pandas provides intuitive commands to save it:

```
df.to_csv('new_dataset.csv')  
df.to_json('new_dataset.json')
```

- When we save JSON and CSV files, all we have to input into those functions is our desired filename with the appropriate file extension.

# Most important Pandas DataFrame operations

- DataFrames possess hundreds of methods and other operations that are crucial to any analysis.
- As a beginner, you should know the operations that:
  - that perform simple transformations of your data and those
  - that provide fundamental statistical analysis on your data.



# Loading dataset

- ▶ We're loading this dataset from a CSV and designating the movie titles to be our index.

```
movies_df = pd.read_csv("movies.csv",  
index_col="title")
```

# Viewing your data

- ▶ The first thing to do when opening a new dataset is print out a few rows to keep as a visual reference. We accomplish this with `.head()`:

```
movies_df.head()
```

- ▶ `.head()` outputs the first five rows of your DataFrame by default, but we could also pass a number as well: `movies_df.head(10)` would output the top ten rows, for example.

- ▶ To see the last five rows use `.tail()` that also accepts a number, and in this case we printing the bottom two rows.:

```
movies_df.tail(2)
```

# Getting info about your data

- ▶ `.info()` should be one of the very first commands you run after loading your data
- ▶ `.info()` provides the essential details about your dataset, such as the number of rows and columns, the number of non-null values, what type of data is in each column, and how much memory your DataFrame is using.

```
movies_df.info()
```

```
movies_df.shape
```

OUT:

```
<class 'pandas.core.frame.DataFrame'>  
Index: 1000 entries, Guardians of the Galaxy to Nine Lives  
Data columns (total 11 columns):  
Rank                1000 non-null int64  
Genre               1000 non-null object  
Description         1000 non-null object  
Director           1000 non-null object  
Actors              1000 non-null object  
Year               1000 non-null int64  
Runtime (Minutes)  1000 non-null int64  
Rating             1000 non-null float64  
Votes              1000 non-null int64  
Revenue (Millions) 872 non-null float64  
Metascore          936 non-null float64  
dtypes: float64(3), int64(4), object(4)  
memory usage: 93.8+ KB
```

OUT:

```
(1000, 11)
```

# Handling duplicates

- ▶ This dataset does not have duplicate rows, but it is always important to verify you aren't aggregating duplicate rows.
- ▶ To demonstrate, let's simply just double up our movies DataFrame by appending it to itself:
- ▶ Using `append()` will return a copy without affecting the original DataFrame. We are capturing this copy in `temp` so we aren't working with the real data.
- ▶ Notice call `.shape` quickly proves our DataFrame rows have doubled.

```
temp_df = movies_df.append(movies_df)
temp_df.shape
```

OUT:

(2000, 11)

Now we can try dropping duplicates:

```
temp_df = temp_df.drop_duplicates()
temp_df.shape
```

OUT:

(1000, 11)

# Handling duplicates

- ▶ Just like `append()`, the `drop_duplicates()` method will also return a copy of your DataFrame, but this time with duplicates removed. Calling `.shape` confirms we're back to the 1000 rows of our original dataset.
- ▶ It's a little verbose to keep assigning DataFrames to the same variable like in this example. For this reason, pandas has the `inplace` keyword argument on many of its methods. Using `inplace=True` will modify the DataFrame object in place:

```
temp_df.drop_duplicates(inplace=True)
```

- ▶ Another important argument for `drop_duplicates()` is `keep`, which has three possible options:
  - ▶ **first**: (default) Drop duplicates except for the first occurrence.
  - ▶ **last**: Drop duplicates except for the last occurrence.
  - ▶ **False**: Drop all duplicates.

# Understanding your variables

- ▶ Using `.describe()` on an entire DataFrame we can get a summary of the distribution of continuous variables:

```
movies_df.describe()
```

OUT:					
	rank	year	runtime	rating	
count	1000.000000	1000.000000	1000.000000	1000.000000	1.00
mean	500.500000	2012.783000	113.172000	6.723200	1.69
std	288.819436	3.205962	18.810908	0.945429	1.88
min	1.000000	2006.000000	66.000000	1.900000	6.10
25%	250.750000	2010.000000	100.000000	6.200000	3.6
50%	500.500000	2014.000000	111.000000	6.800000	1.10
75%	750.250000	2016.000000	123.000000	7.400000	2.3
max	1000.000000	2016.000000	191.000000	9.000000	1.79

- ▶ `.describe()` can also be used on a categorical variable to get the count of rows, unique count of categories, top category, and freq of top category:

```
movies_df['genre'].describe()
```

```
OUT:
count      1000
unique      207
top      Action,Adventure,Sci-Fi
freq         50
Name: genre, dtype: object
```

- ▶ This tells us that the genre column has 207 unique values, the top value is Action/Adventure/Sci-Fi, which shows up 50 times (freq).



# More Examples

```
import pandas as pd
data = [1,2,3,10,20,30]
df = pd.DataFrame(data)
print(df)
```



	0
0	1
1	2
2	3
3	10
4	20
5	30

```
import pandas as pd
data = {'Name' : ['AA', 'BB'], 'Age' : [30,45]}
df = pd.DataFrame(data)
print(df)
```



	Name	Age
0	AA	30
1	BB	45

# More Examples

```
import pandas as pd
data = [{'a': 1, 'b': 2}, {'a': 5, 'b': 10, 'c': 20}]
df = pd.DataFrame(data)
print(df)
```



	a	b	c
0	1	2	NaN
1	5	10	20.0

---

```
import pandas as pd
data = [{'a': 1, 'b': 2}, {'a': 5, 'b': 10, 'c': 20}]
df = pd.DataFrame(data, index=['first', 'second'])
print(df)
```



	a	b	c
first	1	2	NaN
second	5	10	20.0

# More Examples

E.g. This shows how to create a DataFrame with a list of dictionaries, row indices, and column indices.

```
import pandas as pd
data = [{'a': 1, 'b': 2}, {'a': 5, 'b': 10, 'c': 20}]

#With two column indices, values same as dictionary keys
df1 = pd.DataFrame(data, index=['first', 'second'], columns=['a', 'b'])

#With two column indices with one index with other name
df2 = pd.DataFrame(data, index=['first', 'second'], columns=['a', 'b1'])

print(df1)
print('.....')
print(df2)
```

	a	b
first	1	2
second	5	10
.....		
	a	b1
first	1	NaN
second	5	NaN

# More Examples:

## Create a DataFrame from Dict of Series

```
import pandas as pd
d = {'one' : pd.Series([1, 2, 3] , index=['a', 'b', 'c']),
     'two' : pd.Series([1,2, 3, 4], index=['a', 'b', 'c', 'd'])
}
df = pd.DataFrame(d)
print(df)
```

	one	two
a	1.0	1
b	2.0	2
c	3.0	3
d	NaN	4

# More Examples: Column Addition

```
import pandas as pd
d = {'one':pd.Series([1,2,3], index=['a','b','c']),
     'two':pd.Series([1,2,3,4], index=['a','b','c','d'])
}
df = pd.DataFrame(d)
# Adding a new column to an existing DataFrame object
# with column label by passing new series

print("Adding a new column by passing as Series:")
df['three'] = pd.Series([10,20,30],index=['a','b','c'])
print(df)

print("Adding a column using an existing columns in
DataFrame:")
df['four'] = df['one']+df['three']
print(df)
```

Adding a column using Series:

	one	two	three
a	1.0	1	10.0
b	2.0	2	20.0
c	3.0	3	30.0
d	NaN	4	NaN

Adding a column using columns:

	one	two	three	four
a	1.0	1	10.0	11.0
b	2.0	2	20.0	22.0
c	3.0	3	30.0	33.0
d	NaN	4	NaN	NaN

# More Examples: Column Deletion

```
# Using the previous DataFrame, we will delete a column
# using del function
import pandas as pd
d = {'one'      : pd.Series([1, 2, 3],      index=['a', 'b', 'c']),
      'two'      : pd.Series([1, 2, 3, 4],  index=['a', 'b', 'c', 'd']),
      'three'    : pd.Series([10,20,30],    index=['a','b','c'])
    }
df = pd.DataFrame(d)
print ("Our dataframe is:")
print(df)

# using del function
print("Deleting the first column using DEL function:")
del df['one']
print(df)

# using pop function
print("Deleting another column using POP function:")
df.pop('two')
print(df)
```

Our dataframe is:

	one	two	three
a	1.0	1	10.0
b	2.0	2	20.0
c	3.0	3	30.0
d	NaN	4	NaN

Deleting the first column:

	two	three
a	1	10.0
b	2	20.0
c	3	30.0
d	4	NaN

Deleting another column:

a	10.0
b	20.0
c	30.0
d	NaN



# More Examples: Slicing in DataFrames

```
import pandas as pd
d = {'one' : pd.Series([1, 2, 3], index=['a', 'b', 'c']),
     'two' : pd.Series([1, 2, 3, 4], index=['a', 'b', 'c', 'd'])
}
df = pd.DataFrame(d)
print(df[2:4])
```

	one	two
c	3.0	3
d	NaN	4

# More Examples: Addition of rows

```
import pandas as pd
d = {'one' : pd.Series([1, 2, 3], index=['a', 'b', 'c']),
     'two' : pd.Series([1, 2, 3, 4], index=['a', 'b', 'c', 'd'])
}
df = pd.DataFrame(d)
print(df)

df2 = pd.DataFrame([[5,6], [7,8]], columns = ['a', 'b'])
df = df.append(df2 )
print(df)
```

	one	two		
a	1.0	1		
b	2.0	2		
c	3.0	3		
d	NaN	4		

	one	two	a	b
a	1.0	1.0	NaN	NaN
b	2.0	2.0	NaN	NaN
c	3.0	3.0	NaN	NaN
d	NaN	4.0	NaN	NaN
0	NaN	NaN	5.0	6.0
1	NaN	NaN	7.0	8.0

# More Examples: Deletion of rows

```
import pandas as pd
d = {'one':pd.Series([1, 2, 3], index=['a', 'b', 'c']),
     'two':pd.Series([1, 2, 3, 4], index=['a', 'b', 'c', 'd'])
}
df = pd.DataFrame(d)
print(df)

df2 = pd.DataFrame([[5,6], [7,8]], columns = ['a', 'b'])
df = df.append(df2 )
print(df)

df = df.drop(0)
print(df)
```

	one	two
a	1.0	1
b	2.0	2
c	3.0	3
d	NaN	4

	one	two	a	b
a	1.0	1.0	NaN	NaN
b	2.0	2.0	NaN	NaN
c	3.0	3.0	NaN	NaN
d	NaN	4.0	NaN	NaN
0	NaN	NaN	5.0	6.0
1	NaN	NaN	7.0	8.0

	one	two	a	b
a	1.0	1.0	NaN	NaN
b	2.0	2.0	NaN	NaN
c	3.0	3.0	NaN	NaN
d	NaN	4.0	NaN	NaN
1	NaN	NaN	7.0	8.0

# More Examples: Reindexing

```
import pandas as pd
# Creating the first dataframe
df1 = pd.DataFrame({"A": [1, 5, 3, 4, 2],
                    "B": [3, 2, 4, 3, 4],
                    "C": [2, 2, 7, 3, 4],
                    "D": [4, 3, 6, 12, 7]},
                    index = ["A1", "A2", "A3", "A4", "A5"])

# Creating the second dataframe
df2 = pd.DataFrame({"A": [10, 11, 7, 8, 5],
                    "B": [21, 5, 32, 4, 6],
                    "C": [11, 21, 23, 7, 9],
                    "D": [1, 5, 3, 8, 6]},
                    index = ["A1", "A3", "A4", "A7", "A8"])

# Print the first dataframe
print(df1)
print(df2)
# find matching indexes
df1.reindex_like(df2)
```

- Pandas `dataframe.reindex_like()` function return an object with matching indices to myself.
- Any non-matching indexes are filled with NaN values.

Out[72]:

	A	B	C	D
A1	1.0	3.0	2.0	4.0
A3	3.0	4.0	7.0	6.0
A4	4.0	3.0	3.0	12.0
A7	NaN	NaN	NaN	NaN
A8	NaN	NaN	NaN	NaN

# More Examples:

## Concatenating Objects (Data Frames)

```
import pandas as pd
df1 = pd.DataFrame({'Name': ['A', 'B'], 'SSN': [10, 20], 'marks': [90, 95] })
df2 = pd.DataFrame({'Name': ['B', 'C'], 'SSN': [25, 30], 'marks': [80, 97] })
df3 = pd.concat([df1, df2])
df3
```

# Handling categorical data

- ▶ There are many data that are repetitive for example gender , country , and codes are always repetitive .
- ▶ Categorical variables can take on only a limited
- ▶ The categorical data type is useful in the following cases –
- ▶ A string variable consisting of only a few different values. Converting such a string variable to a categorical variable will save some memory.
- ▶ The lexical order of a variable is not the same as the logical order (“one”, “two”, “three”).
  - ▶ By converting to a categorical and specifying an order on the categories, sorting and min/max will use the logical order instead of the lexical order.
- ▶ As a signal to other python libraries that this column should be treated as a categorical variable (e.g. to use suitable statistical methods or plot types).

# Examples

```
import pandas as pd
cat = pd.Categorical(['a', 'b', 'c', 'a', 'b', 'c'])
print(cat)
```

```
import pandas as pd
import numpy as np
cat = pd.Categorical(["a", "c", "c", np.nan], categories=["b", "a", "c"])
df = pd.DataFrame({"cat": cat, "s": ["a", "c", "c", np.nan]})
print(df.describe())
print(df["cat"].describe())
```

# Reading data from a SQL database

- ▶ If you're working with data from a SQL database you need to first establish a connection using an appropriate Python library, then pass a query to pandas. Here we'll use SQLite to demonstrate.
- ▶ First, we need pysqlite3 installed, so run this command in your terminal:
  - ▶ `pip install pysqlite3`
  - ▶ Or run this cell if you're in a notebook: `!pip install pysqlite3`
- ▶ `sqlite3` is used to create a connection to a database which we can then use to generate a DataFrame through a `SELECT` query.
  - ▶ So first we'll make a connection to a SQLite database file:

```
import sqlite3  
con = sqlite3.connect("database.db")
```

- ▶ In this SQLite database we have a table called `purchases`, and our index is in a column called `"index"`.
- ▶ By passing a `SELECT` query and our `con`, we can read from the `purchases` table:

```
df = pd.read_sql_query("SELECT * FROM purchases", con)
```



# Reading data from a SQL database

- ▶ In this SQLite database we have a table called purchases, and our index is in a column called "index".
- ▶ By passing a SELECT query and our con, we can read from the purchases table:

```
df = pd.read_sql_query("SELECT * FROM purchases", con)
```

OUT:

	index	apples	oranges
0	June	3	0
1	Robert	2	3
2	Lily	0	7
3	David	1	2

- ▶ Just like with CSVs, we could pass index\_col='index', but we can also set an index after-the-fact:
  - ▶ In fact, we could use set\_index() on any DataFrame using any column at any time. Indexing Series and DataFrames is a very common task, and the different ways of doing it is worth remembering.

```
df = df.set_index('index')
```

OUT:

	apples	oranges
index		
June	3	0
Robert	2	3
Lily	0	7
David	1	2

# References

## pandas documentation

- ▶ <https://pandas.pydata.org/pandas-docs/stable/index.html>
- ▶ pandas: Input/output
  - ▶ <https://pandas.pydata.org/pandas-docs/stable/reference/io.html>
- ▶ pandas: DataFrame
  - ▶ <https://pandas.pydata.org/pandas-docs/stable/reference/frame.html>
- ▶ pandas: Series
  - ▶ <https://pandas.pydata.org/pandas-docs/stable/reference/series.html>
- ▶ pandas: Plotting
  - ▶ <https://pandas.pydata.org/pandas-docs/stable/reference/plotting.html>

# Python For Data Science Cheat Sheet

## Pandas Basics

Learn Python for Data Science Interactively at [www.DataCamp.com](https://www.datacamp.com)



### Pandas

The Pandas library is built on NumPy and provides easy-to-use data structures and data analysis tools for the Python programming language.



Use the following import convention:

```
>>> import pandas as pd
```

### Pandas Data Structures

#### Series

A one-dimensional labeled array capable of holding any data type

A	3
B	-5
C	7
D	4

```
>>> s = pd.Series([3, -5, 7, 4], index=['a', 'b', 'c', 'd'])
```

#### DataFrame

Columns			
	Country	Capital	Population
1	Belgium	Brussels	11190846
2	India	New Delhi	1303171035
3	Brazil	Brasilia	207847528

A two-dimensional labeled data structure with columns of potentially different types

```
>>> data = {'Country': ['Belgium', 'India', 'Brazil'],
           'Capital': ['Brussels', 'New Delhi', 'Brasilia'],
           'Population': [11190846, 1303171035, 207847528]}
```

```
>>> df = pd.DataFrame(data,
                      columns=['Country', 'Capital', 'Population'])
```

### I/O

#### Read and Write to CSV

```
>>> pd.read_csv('file.csv', header=None, nrows=5)
>>> pd.to_csv('myDataFrame.csv')
```

#### Read and Write to Excel

```
>>> pd.read_excel('file.xlsx')
>>> pd.to_excel('dir/myDataFrame.xlsx', sheet_name='Sheet1')
```

Read multiple sheets from the same file

```
>>> xlsx = pd.ExcelFile('file.xls')
>>> df = pd.read_excel(xlsx, 'Sheet1')
```

### Asking For Help

```
>>> help(pd.Series.loc)
```

### Selection

Also see NumPy Arrays

#### Getting

```
>>> s['b']
-5
Get one element

>>> df[1:]
   Country  Capital  Population
1  India   New Delhi  1303171035
2  Brazil  Brasilia  207847528
Get subset of a DataFrame
```

#### Selecting, Boolean Indexing & Setting

##### By Position

```
>>> df.iloc([0], [0])
'Belgium'
Select single value by row & column

>>> df.iat([0], [0])
'Belgium'
```

##### By Label

```
>>> df.loc([0], ['Country'])
'Belgium'
Select single value by row & column labels

>>> df.at([0], ['Country'])
'Belgium'
```

##### By Label/Position

```
>>> df.ix[2]
Country      Brazil
Capital      Brasilia
Population    207847528
Select single row of subset of rows
```

```
>>> df.ix[:, 'Capital']
0    Brussels
1    New Delhi
2    Brasilia
Select a single column of subset of columns
```

```
>>> df.ix[1, 'Capital']
'New Delhi'
Select rows and columns
```

##### Boolean Indexing

```
>>> s[~(s > 1)]
>>> s[(s < -1) | (s > 2)]
>>> df[df['Population'] > 1200000000]
Use filter to adjust DataFrame
```

##### Setting

```
>>> s['a'] = 6
Set index a of Series s to 6
```

### Dropping

```
>>> s.drop(['a', 'c'])
Drop values from rows (axis=0)
>>> df.drop('Country', axis=1)
Drop values from columns (axis=1)
```

### Sort & Rank

```
>>> df.sort_index()
Sort by labels along an axis
>>> df.sort_values(by='Country')
Sort by the values along an axis
>>> df.rank()
Assign ranks to entries
```

### Retrieving Series/DataFrame Information

#### Basic Information

```
>>> df.shape
(rows, columns)
>>> df.index
Describe index
>>> df.columns
Describe DataFrame columns
>>> df.info()
Info on DataFrame
>>> df.count()
Number of non-NA values
```

#### Summary

```
>>> df.sum()
Sum of values
>>> df.cumsum()
Cumulative sum of values
>>> df.min()/df.max()
Minimum/maximum values
>>> df.idxmin()/df.idxmax()
Minimum/Maximum index value
>>> df.describe()
Summary statistics
>>> df.mean()
Mean of values
>>> df.median()
Median of values
```

### Applying Functions

```
>>> f = lambda x: x*2
>>> df.apply(f)
Apply function
>>> df.applymap(f)
Apply function element-wise
```

### Data Alignment

#### Internal Data Alignment

NA values are introduced in the indices that don't overlap:

```
>>> s3 = pd.Series([7, -2, 3], index=['a', 'c', 'd'])
>>> s + s3
a    10.0
b     NaN
c     5.0
d     7.0
```

#### Arithmetic Operations with Fill Methods

You can also do the internal data alignment yourself with the help of the fill methods:

```
>>> s.add(s3, fill_value=0)
a    10.0
b    -5.0
c     5.0
d     7.0
>>> s.sub(s3, fill_value=2)
>>> s.div(s3, fill_value=4)
>>> s.mul(s3, fill_value=3)
```

DataCamp

Learn Python for Data Science Interactively





# Data Wrangling

## with pandas Cheat Sheet

<http://pandas.pydata.org>

### Syntax – Creating DataFrames

	a	b	c
1	4	7	10
2	5	8	11
3	6	9	12

```
df = pd.DataFrame(  
    {"a" : [4, 5, 6],  
     "b" : [7, 8, 9],  
     "c" : [10, 11, 12]},  
    index = [1, 2, 3])  
Specify values for each column.
```

```
df = pd.DataFrame(  
    [[4, 7, 10],  
     [5, 8, 11],  
     [6, 9, 12]],  
    index=[1, 2, 3],  
    columns=['a', 'b', 'c'])  
Specify values for each row.
```

		a	b	c
n	1	4	7	10
d	2	5	8	11
e	2	6	9	12


```
df = pd.DataFrame(  
    {"a" : [4, 5, 6],  
     "b" : [7, 8, 9],  
     "c" : [10, 11, 12]},  
    index = pd.MultiIndex.from_tuples(  
        [('d', 1), ('d', 2), ('e', 2)],  
        names=['n', 'v']))  
Create DataFrame with a MultiIndex
```

### Method Chaining

Most pandas methods return a DataFrame so that another pandas method can be applied to the result. This improves readability of code.

```
df = (pd.melt(df)  
      .rename(columns={  
          'variable' : 'var',  
          'value' : 'val'})  
      .query('val >= 200'))
```

### Tidy Data – A foundation for wrangling in pandas



In a tidy data set:

Each **variable** is saved in its own **column**



Each **observation** is saved in its own **row**

Tidy data complements pandas's **vectorized operations**. pandas will automatically preserve observations as you manipulate variables. No other format works as intuitively with pandas.




M \* A


### Reshaping Data – Change the layout of a data set



```
pd.melt(df)  
Gather columns into rows.
```



```
df.pivot(columns='var', values='val')  
Spread rows into columns.
```




```
pd.concat([df1, df2])  
Append rows of DataFrames
```



```
pd.concat([df1, df2], axis=1)  
Append columns of DataFrames
```

```
df.sort_values('mpg')  
Order rows by values of a column (low to high).  
  
df.sort_values('mpg', ascending=False)  
Order rows by values of a column (high to low).  
  
df.rename(columns = {'y': 'year'})  
Rename the columns of a DataFrame  
  
df.sort_index()  
Sort the index of a DataFrame  
  
df.reset_index()  
Reset index of DataFrame to row numbers, moving index to columns.  
  
df.drop(columns=['Length', 'Height'])  
Drop columns from DataFrame
```

### Subset Observations (Rows)



```
df[df.Length > 7]
```

Extract rows that meet logical criteria.

```
df.drop_duplicates()
```

Remove duplicate rows (only considers columns).

```
df.head(n)
```

Select first n rows.

```
df.tail(n)
```

Select last n rows.

```
df.sample(frac=0.5)
```

Randomly select fraction of rows.

```
df.sample(n=10)
```

Randomly select n rows.

```
df.iloc[10:20]
```

Select rows by position.


```
df.nlargest(n, 'value')
```

Select and order top n entries.

```
df.nsmallest(n, 'value')
```

Select and order bottom n entries.

### Subset Variables (Columns)



```
df[['width', 'length', 'species']]
```

Select multiple columns with specific names.

```
df['width'] or df.width
```

Select single column with specific name.

```
df.filter(regex='regex')
```

Select columns whose name matches regular expression *regex*.

#### regex (Regular Expressions) Examples

regex	Matches
'\.'	Matches strings containing a period '.'
'Length\$'	Matches strings ending with word 'Length'
'^Sepal'	Matches strings beginning with the word 'Sepal'
'^x[1-5]\$'	Matches strings beginning with 'x' and ending with 1,2,3,4,5
'^(?!Species\$).*\$'	Matches strings except the string 'Species'

```
df.loc[:, 'x2': 'x4']
```

Select all columns between x2 and x4 (inclusive).

```
df.iloc[:, [1, 2, 5]]
```

Select columns in positions 1, 2 and 5 (first column is 0).

```
df.loc[df['a'] > 10, ['a', 'c']]
```

Select rows meeting logical condition, and only the specific columns.

Logic in Python (and pandas)			
<	Less than	!=	Not equal to
>	Greater than	df.column.isin(values)	Group membership
==	Equals	pd.isnull(obj)	Is NaN
<=	Less than or equals	pd.notnull(obj)	Is not NaN
>=	Greater than or equals	&,  , ~, ^, df.any(), df.all()	Logical and, or, not, xor, any, all



## Summarize Data

**df['Length'].value\_counts()**

Count number of rows with each unique value of variable

**len(df)**


# of rows in DataFrame.

**len(df['w'].unique())**

# of distinct values in a column.

**df.describe()**

Basic descriptive statistics for each column (or GroupBy)



pandas provides a large set of **summary functions** that operate on different kinds of pandas objects (DataFrame columns, Series, GroupBy, Expanding and Rolling (see below)) and produce single values for each of the groups. When applied to a DataFrame, the result is returned as a pandas Series for each column. Examples:

**sum()**

Sum values of each object.

**count()**

Count non-NA/null values of each object.

**median()**

Median value of each object.

**quantile([0.25,0.75])**

Quantiles of each object.

**apply(function)**

Apply function to each object.

**min()**

Minimum value in each object.

**max()**

Maximum value in each object.

**mean()**

Mean value of each object.


**var()**

Variance of each object.

**std()**

Standard deviation of each object.

## Group Data



**df.groupby(by="col")**

Return a GroupBy object, grouped by values in column named "col".

**df.groupby(level="ind")**

Return a GroupBy object, grouped by values in index level named "ind".

All of the summary functions listed above can be applied to a group. Additional GroupBy functions:

**size()**

Size of each group.

**agg(function)**

Aggregate group using function.

## Windows

**df.expanding()**

Return an Expanding object allowing summary functions to be applied cumulatively.

**df.rolling(n)**

Return a Rolling object allowing summary functions to be applied to windows of length n.

## Handling Missing Data


**df=df.dropna()**

Drop rows with any column having NA/null data.

**df=df.fillna(value)**

Replace all NA/null data with value.

## Make New Variables



**df=df.assign(Area=lambda df: df.Length\*df.Height)**


Compute and append one or more new columns.

**df['Volume'] = df.Length\*df.Height\*df.Depth**

Add single column.

**pd.qcut(df.col, n, labels=False)**

Bin column into n buckets.



pandas provides a large set of **vector functions** that operate on all columns of a DataFrame or a single selected column (a pandas Series). These functions produce vectors of values for each of the columns, or a single Series for the individual Series. Examples:

**max(axis=1)**

Element-wise max.

**min(axis=1)**

Element-wise min.

**clip(lower=-10,upper=10) abs()**

Trim values at input thresholds Absolute value.

The examples below can also be applied to groups. In this case, the function is applied on a per-group basis, and the returned vectors are of the length of the original DataFrame.

**shift(1)**

Copy with values shifted by 1.

**rank(method='dense')**

Ranks with no gaps.

**rank(method='min')**

Ranks. Ties get min rank.

**rank(pct=True)**

Ranks rescaled to interval [0, 1].

**rank(method='first')**

Ranks. Ties go to first value.

**shift(-1)**

Copy with values lagged by 1.

**cumsum()**

Cumulative sum.

**cummax()**

Cumulative max.

**cummin()**

Cumulative min.


**cumprod()**

Cumulative product.

## Plotting


**df.plot.hist()**

Histogram for each column



**df.plot.scatter(x='w',y='h')**

Scatter chart using pairs of points



## Combine Data Sets

**adf**

x1	x2
A	1
B	2
C	3

**bdf**

x1	x3
A	T
B	F
D	T



Standard Joins

x1	x2	x3
A	1	T
B	2	F
C	3	NaN

**pd.merge(adf, bdf, how='left', on='x1')**  
Join matching rows from bdf to adf.

x1	x2	x3
A	1.0	T
B	2.0	F
D	NaN	T

**pd.merge(adf, bdf, how='right', on='x1')**  
Join matching rows from adf to bdf.

x1	x2	x3
A	1	T
B	2	F

**pd.merge(adf, bdf, how='inner', on='x1')**  
Join data. Retain only rows in both sets.

x1	x2	x3
A	1	T
B	2	F
C	3	NaN
D	NaN	T

**pd.merge(adf, bdf, how='outer', on='x1')**  
Join data. Retain all values, all rows.

Filtering Joins

x1	x2
A	1
B	2

**adf[adf.x1.isin(bdf.x1)]**  
All rows in adf that have a match in bdf.

x1	x2
C	3

**adf[~adf.x1.isin(bdf.x1)]**  
All rows in adf that do not have a match in bdf.

**ydf**

x1	x2
A	1
B	2
C	3



**zdf**

x1	x2
B	2
C	3
D	4



Set-like Operations

x1	x2
B	2
C	3

**pd.merge(ydf, zdf)**  
Rows that appear in both ydf and zdf (Intersection).

x1	x2
A	1
B	2
C	3
D	4

**pd.merge(ydf, zdf, how='outer')**  
Rows that appear in either or both ydf and zdf (Union).

x1	x2
A	1

**pd.merge(ydf, zdf, how='outer', indicator=True).query('\_merge == "left\_only")**  
**.drop(['\_merge'],axis=1)**  
Rows that appear in ydf but not zdf (Setdiff).

## Quick start


```
import numpy as np
import matplotlib as mpl
import matplotlib.pyplot as plt
```

```
X = np.linspace(0, 2*np.pi, 100)
Y = np.cos(X)
```

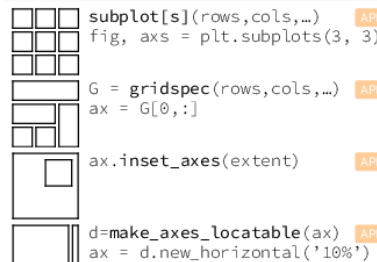
```
fig, ax = plt.subplots()
ax.plot(X, Y, color='green')
```

```
fig.savefig("figure.pdf")
plt.show()
```

## Anatomy of a figure



## Subplots layout



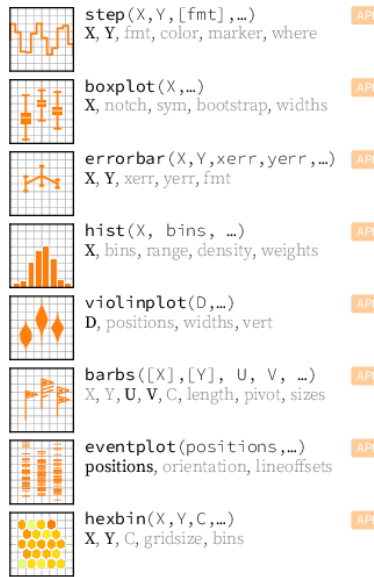
## Getting help

[matplotlib.org](https://matplotlib.org)  
[github.com/matplotlib/matplotlib/issues](https://github.com/matplotlib/matplotlib/issues)  
[discourse.matplotlib.org](https://discourse.matplotlib.org)  
[stackoverflow.com/questions/tagged/matplotlib](https://stackoverflow.com/questions/tagged/matplotlib)  
<https://gitter.im/matplotlib/matplotlib>  
[twitter.com/matplotlib](https://twitter.com/matplotlib)  
[matplotlib users mailing list](mailto:matplotlib-users@googlegroups.com)


## Basic plots




## Advanced plots




## Scales




## Projections



## Lines



## Markers



## Colors

